

Some other Possible Process State

- A process could be **Waiting** for something to happen
 ↳ Example: parent process waiting for child process to terminate
 through a `wait()` system call
- A process could be **Ready** for the OS to cause it to run
 ↳ Example: After a child process terminates, the parent process may not be able to run right away because some other process is running.
- Running, Waiting, Ready → to be allowed to run by the OS
 on CPU ↓ for something to happen ↓

Process Management

- OS job of managing processes that are waiting, running or ready
- What should the OS do when a process does something that takes a long time?

An OS designer Question
 ↳ eg anything that involves the hard disk access
 ↳ file read/write operation, page fault
 ↓
 does give us perspective

①. Solution → do nothing.

→ Problem ⇒ processor is idle billions of cycles

hd seconds CPU nseconds
 $\frac{1s}{1ns} = 10^9$

⇒ the processor could have executed billions of instructions instead during that time

② OS should try to maximize processor **utilization**

utilization ⇒ fraction of the time that the processor is busy → get to 100%

(...)

busy → get to 100%
How?

- OS could change status of that process to "waiting" and make another process "Running"

Question: Which other process?

↳ Determined by the process scheduler

Process Scheduler

- The part of the OS that manages the sharing of CPU time among processes.
- Possible considerations that the scheduler could use in making scheduling decisions.

↳ Minimize average program execution time

$$T_{ave} = \frac{T_{prog1} + T_{prog2} + T_{prog3}}{3}$$

↓
make scheduling decisions based on well being of all programs

BUT → no guarantee of fairness

↳ Fairness to all the programs in execution

↳ Do not sacrifice one or two processes in the interest of the average value

Process Scheduling Policies

Idea 1: let the currently running process continue to do so until it does something that involves a long time

↳ disk access page fault

Non-preemptive policy

↳ then switch to a Ready Process

very dangerous { ↳ Problem → what happens if the currently running process is creating an infinite

↳ debugging your program

while(1);

↳ No other process would ever get to run on the OS with Idea 1

↳ NOT used for Linux/Unix systems but used for some embedded systems.

Idea 2 → Preemptive Policy

⇒ one where the OS "preempts" the running process from the CPU even though, the CPU is not waiting for something.

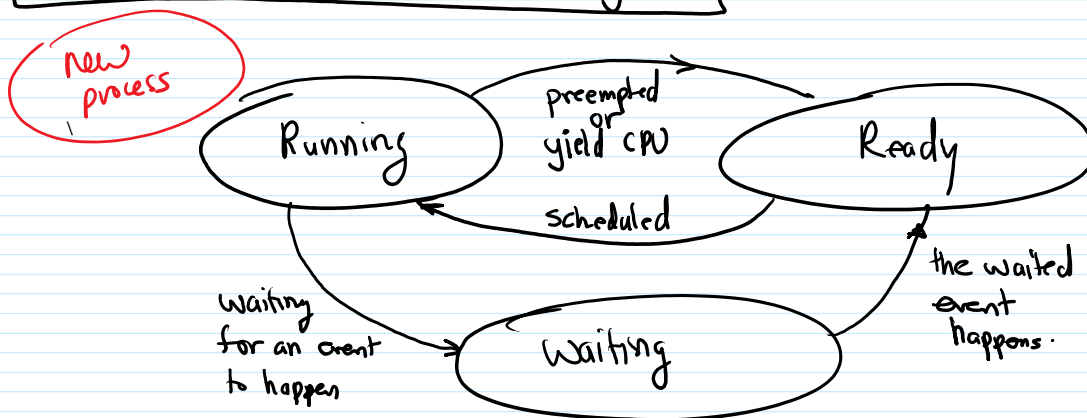
for fairness or benefit of other processes

How? When do I exit?

↓
Idea ⇒ give a maximum amount of CPU time before preempting the process.

CPU time slice ⇒ maximum amount of time allotted to a process before preempting it from the CPU

Process State Transition Diagram



Context Switch


- The event of going from Waiting to Running

- The event of going from Waiting to Running $P1 \rightarrow P4$
- The context switch takes time

contents of PC \rightarrow save HW state of previously running process
 registers, IR etc \rightarrow restore HW state of newly scheduled process

- The amount of time context switch takes determines
CPU slice time

\rightarrow CPU slice time \gg context switch time



\rightarrow generally CPU slice time \sim 1 second.

Non-preemptive Scheduling Policies


① First Come First Serve (FCFS)

Idea: maintain a queue of ready processes.

Queue \Rightarrow a data structure with 2 operations

Insert \rightarrow add new process to the back of the Queue

Delete \rightarrow remove the process from the front of the Queue



\rightarrow schedule next process from the front of the Ready Q

Problem FCFS not a good policy. \Rightarrow doesn't do anything about processes that have small time requirements


② SPN - shortest process next

- The policy which results in the lowest possible average program execution time
 - Schedule next process that which requires the least CPU time
- Problem** → not very practical
→ looking into a future! ⇒ need a good estimate

Preemptive Scheduling Policies

① RR - Round Robin

↳ maintain a FCFS ReadyQ



- much more fair than non-preemptive policies
- As before, when the currently running process is preempted, schedule the process from the front of the ReadyQ
- BUT**
- Insert previous process to the back